

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (accuracy_score, confusion_matrix,
                             classification_report, roc_auc_score,
                             precision_recall_curve,
                             average_precision_score)
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler
from imblearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
import warnings
warnings.filterwarnings("ignore")

df = pd.read_csv(r'C:\Users\Rohit
Kadam\Downloads\supply_chain_data.csv')
df.head()

```

| | Product type | SKU | Price | Availability | Number of products sold |
|---|--------------|------|-----------|--------------|-------------------------|
| 0 | hairecare | SKU0 | 69.808006 | 55 | 802 |
| 1 | skincare | SKU1 | 14.843523 | 95 | 736 |
| 2 | hairecare | SKU2 | 11.319683 | 34 | 8 |
| 3 | skincare | SKU3 | 61.163343 | 68 | 83 |
| 4 | skincare | SKU4 | 4.805496 | 26 | 871 |

| | Revenue generated | Customer demographics | Stock levels | Lead times |
|---|-------------------|-----------------------|--------------|------------|
| 0 | 8661.996792 | Non-binary | 58 | 7 |
| 1 | 7460.900065 | Female | 53 | 30 |
| 2 | 9577.749626 | Unknown | 1 | 10 |
| 3 | 7766.836426 | Non-binary | 23 | 13 |
| 4 | 2686.505152 | Non-binary | 5 | 3 |

| | Order quantities | ... | Location | Lead time | Production volumes | \ |
|---|------------------|-----|----------|-----------|--------------------|---|
| 0 | 96 | ... | Mumbai | 29 | 215 | |
| 1 | 37 | ... | Mumbai | 23 | 517 | |
| 2 | 88 | ... | Mumbai | 12 | 971 | |
| 3 | 59 | ... | Kolkata | 24 | 937 | |
| 4 | 56 | ... | Delhi | 5 | 414 | |

| | Manufacturing lead time | Manufacturing costs | Inspection results | \ |
|---|-------------------------|---------------------|--------------------|---|
| 0 | 29 | 46.279879 | Pending | |
| 1 | 30 | 33.616769 | Pending | |
| 2 | 27 | 30.688019 | Pending | |
| 3 | 18 | 35.624741 | Fail | |
| 4 | 3 | 92.065161 | Fail | |

| | Defect rates | Transportation modes | Routes | Costs |
|---|--------------|----------------------|---------|------------|
| 0 | 0.226410 | Road | Route B | 187.752075 |
| 1 | 4.854068 | Road | Route B | 503.065579 |
| 2 | 4.580593 | Air | Route C | 141.920282 |
| 3 | 4.746649 | Rail | Route A | 254.776159 |
| 4 | 3.145580 | Air | Route A | 923.440632 |

[5 rows x 24 columns]

df.shape

(100, 24)

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 100 entries, 0 to 99

Data columns (total 24 columns):

| # | Column | Non-Null Count | Dtype |
|-----|-------------------------|----------------|---------|
| --- | ----- | ----- | ----- |
| 0 | Product type | 100 non-null | object |
| 1 | SKU | 100 non-null | object |
| 2 | Price | 100 non-null | float64 |
| 3 | Availability | 100 non-null | int64 |
| 4 | Number of products sold | 100 non-null | int64 |
| 5 | Revenue generated | 100 non-null | float64 |
| 6 | Customer demographics | 100 non-null | object |
| 7 | Stock levels | 100 non-null | int64 |
| 8 | Lead times | 100 non-null | int64 |
| 9 | Order quantities | 100 non-null | int64 |
| 10 | Shipping times | 100 non-null | int64 |
| 11 | Shipping carriers | 100 non-null | object |
| 12 | Shipping costs | 100 non-null | float64 |
| 13 | Supplier name | 100 non-null | object |
| 14 | Location | 100 non-null | object |
| 15 | Lead time | 100 non-null | int64 |

16 Production volumes 100 non-null int64
17 Manufacturing lead time 100 non-null int64
18 Manufacturing costs 100 non-null float64
19 Inspection results 100 non-null object
20 Defect rates 100 non-null float64
21 Transportation modes 100 non-null object
22 Routes 100 non-null object
23 Costs 100 non-null float64
dtypes: float64(6), int64(9), object(9)
memory usage: 18.9+ KB

df.head()

Product type SKU Price Availability Number of products sold
0 haircare SKU0 69.808006 55 802
1 skincare SKU1 14.843523 95 736
2 haircare SKU2 11.319683 34 8
3 skincare SKU3 61.163343 68 83
4 skincare SKU4 4.805496 26 871

Revenue generated Customer demographics Stock levels Lead
times \
0 8661.996792 Non-binary 58 7
1 7460.900065 Female 53 30
2 9577.749626 Unknown 1 10
3 7766.836426 Non-binary 23 13
4 2686.505152 Non-binary 5 3

Order quantities ... Location Lead time Production volumes \
0 96 ... Mumbai 29 215
1 37 ... Mumbai 23 517
2 88 ... Mumbai 12 971
3 59 ... Kolkata 24 937
4 56 ... Delhi 5 414

Manufacturing lead time Manufacturing costs Inspection results \
0 29 46.279879 Pending
1 30 33.616769 Pending
2 27 30.688019 Pending
3 18 35.624741 Fail


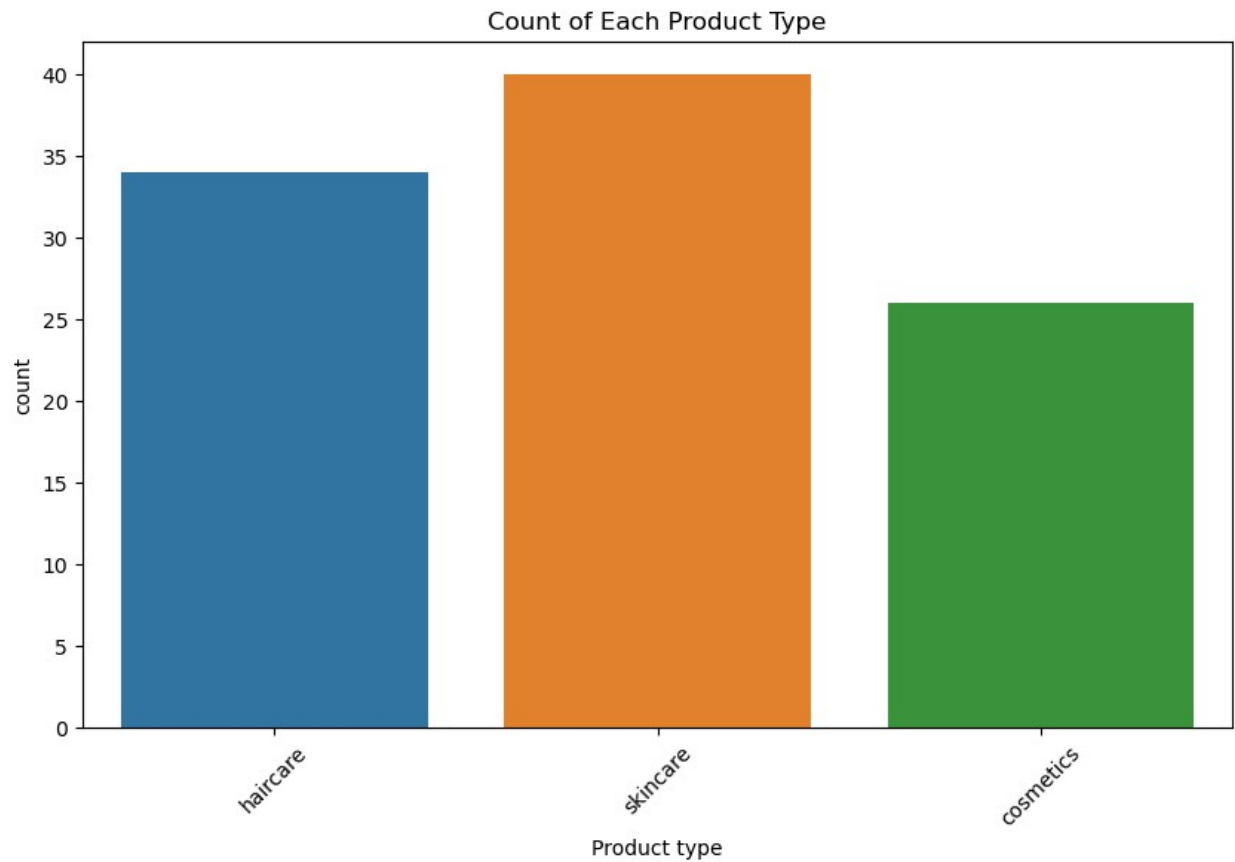
| 4 | | 3 | 92.065161 | | Fail |
|---|--------------|----------------------|-----------|------------|------|
| | Defect rates | Transportation modes | Routes | Costs | |
| 0 | 0.226410 | Road | Route B | 187.752075 | |
| 1 | 4.854068 | Road | Route B | 503.065579 | |
| 2 | 4.580593 | Air | Route C | 141.920282 | |
| 3 | 4.746649 | Rail | Route A | 254.776159 | |
| 4 | 3.145580 | Air | Route A | 923.440632 | |

[5 rows x 24 columns]


```
plt.figure(figsize=(10,6))
sns.countplot(data=df, x='Product type')
plt.title('Count of Each Product Type')
plt.xticks(rotation=45)
plt.show()
```

```
plt.figure(figsize=(12, 6))
sns.boxplot(data=df,x='Product type',y='Price')
plt.title('Price Distribution by Product Type')
plt.show()
```

```
plt.figure(figsize=(12, 6))
sns.scatterplot(data=df,
                x='Number of products sold',
                y='Revenue generated',
                hue='Product type',
                size='Price',
                palette='Set2', # optional: you can change color
                sizes=(40, 400))
```



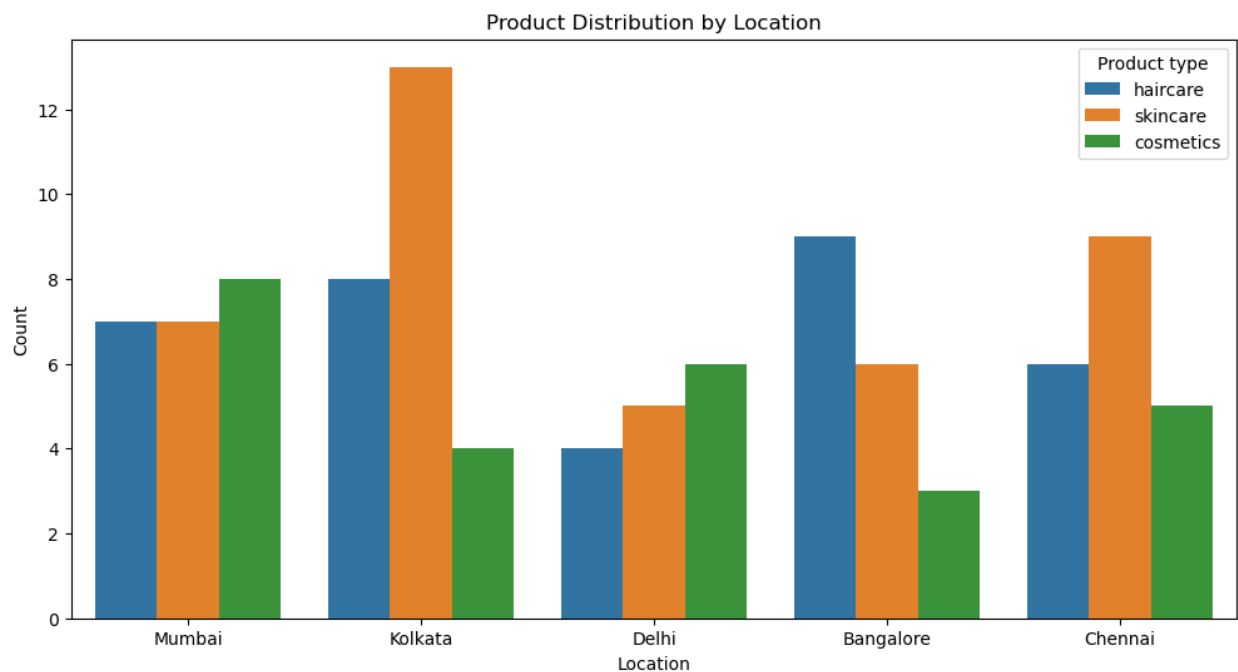
<Axes: xlabel='Number of products sold', ylabel='Revenue generated'>



2. Location and Supplier Analysis

Location-wise product distribution

```
plt.figure(figsize=(12,6))
sns.countplot(data=df,x='Location',hue='Product type')
plt.title('Product Distribution by Location')
plt.ylabel('Count')
plt.show()
```



```
# Supplier performance metrics
supplier_metrics = df.groupby('Supplier name').agg({
    'Lead times': 'mean',
    'Manufacturing costs': 'mean',
    'Defect rates': 'mean'
}).reset_index()
supplier_metrics
```

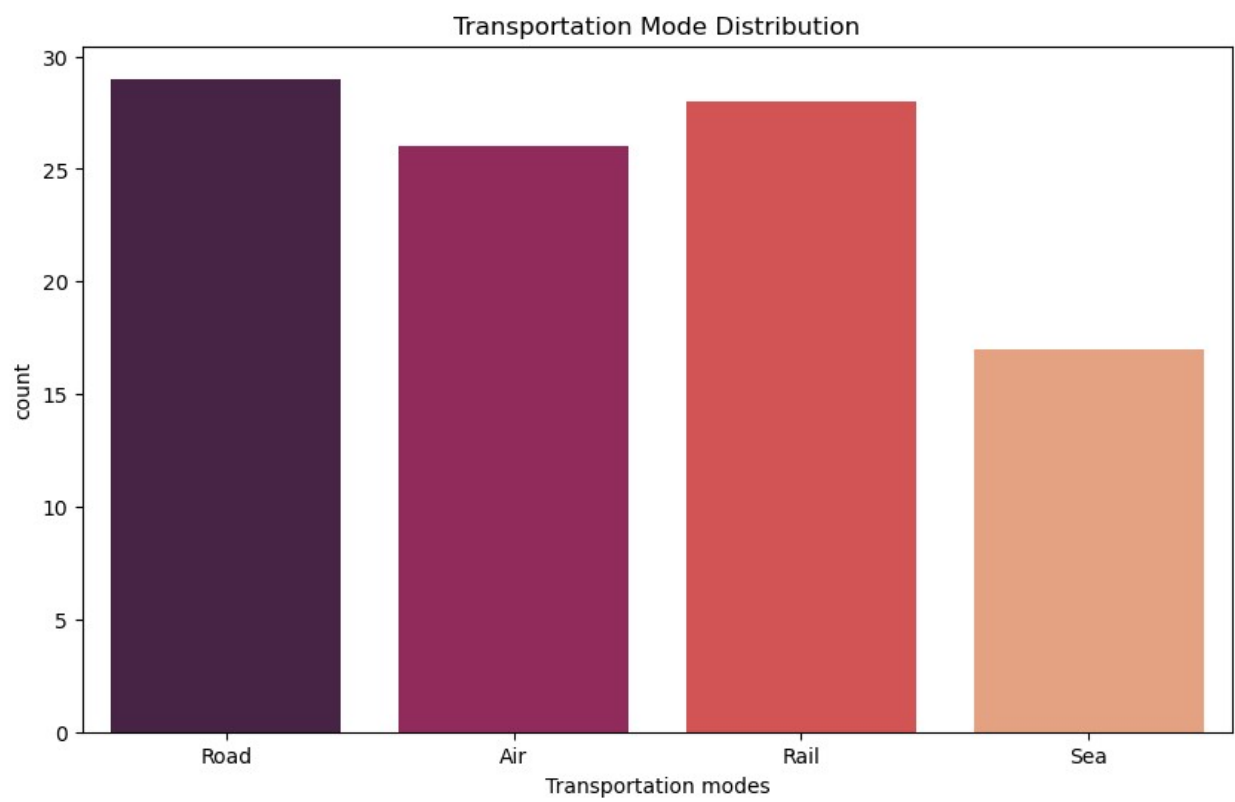

| | Supplier name | Lead times | Manufacturing costs | Defect rates |
|---|---------------|------------|---------------------|--------------|
| 0 | Supplier 1 | 16.777778 | 45.254027 | 1.803630 |
| 1 | Supplier 2 | 16.227273 | 41.622514 | 2.362750 |
| 2 | Supplier 3 | 14.333333 | 43.634121 | 2.465786 |
| 3 | Supplier 4 | 17.000000 | 62.709727 | 2.337397 |
| 4 | Supplier 5 | 14.722222 | 44.768243 | 2.665408 |

```
fig = px.bar(supplier_metrics, x='Supplier name', y=['Lead times',
'Manufacturing costs', 'Defect rates'],
            title='Supplier Performance Metrics', barmode='group')
fig.show()
```

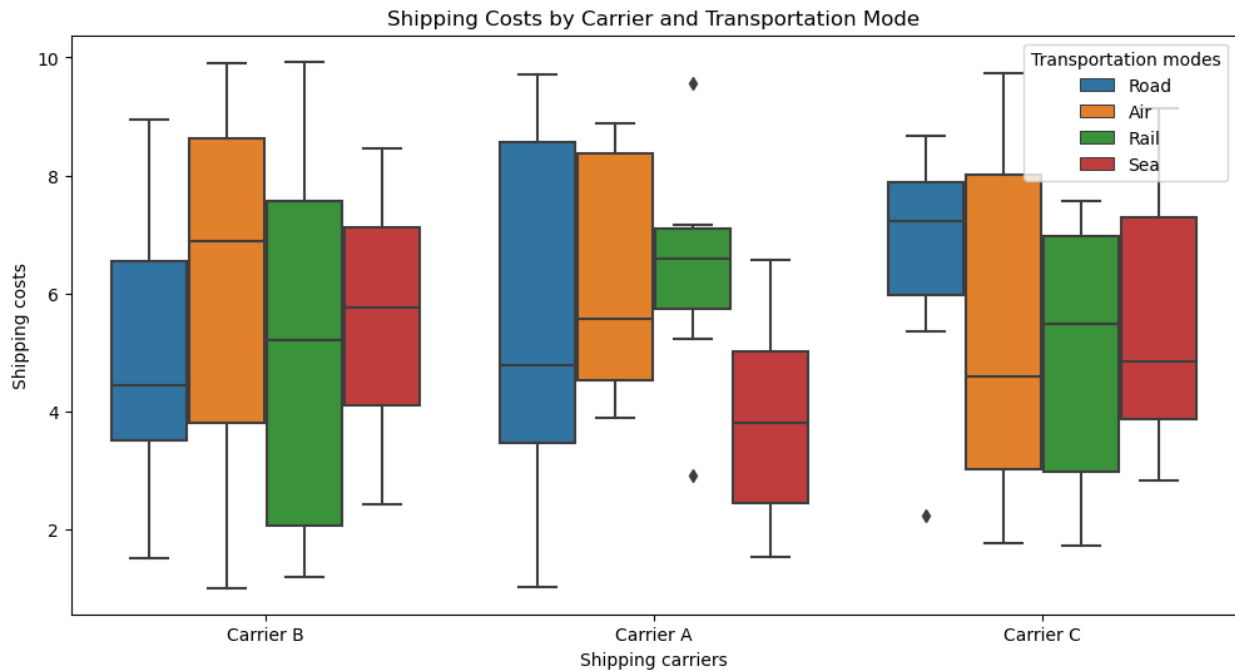


3. Transportation and Logistics Analysis

```
# Transportation mode analysis
plt.figure(figsize=(10, 6))
sns.countplot(x='Transportation modes', data=df, palette="rocket")
plt.title('Transportation Mode Distribution')
plt.show()
plt.figure(figsize=(10, 6))
sns.countplot(x='Transportation modes', data=df, palette="rocket")
plt.title('Transportation Mode Distribution')
plt.show()
```




```
# Shipping costs by carrier
plt.figure(figsize=(12, 6))
sns.boxplot(x='Shipping carriers', y='Shipping costs',
hue='Transportation modes', data=df)
plt.title('Shipping Costs by Carrier and Transportation Mode')
plt.show()
```

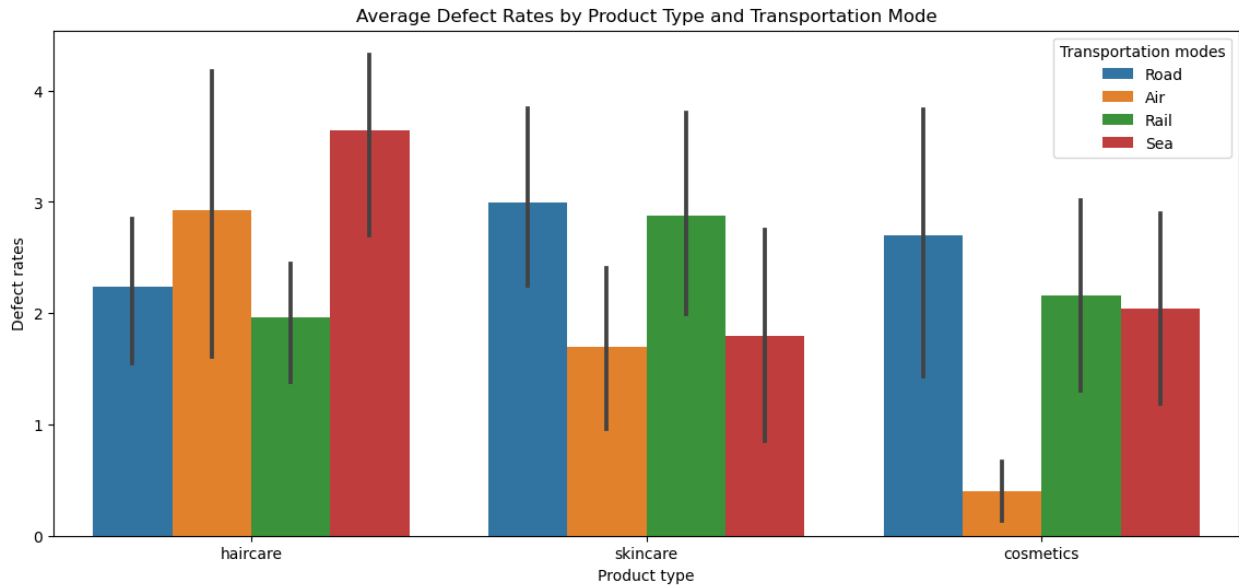


```
# Lead time vs Manufacturing costs
px.scatter(df, x='Lead times', y='Manufacturing costs',
color='Product type', size='Production volumes',
title='Lead Time vs Manufacturing Costs')
```

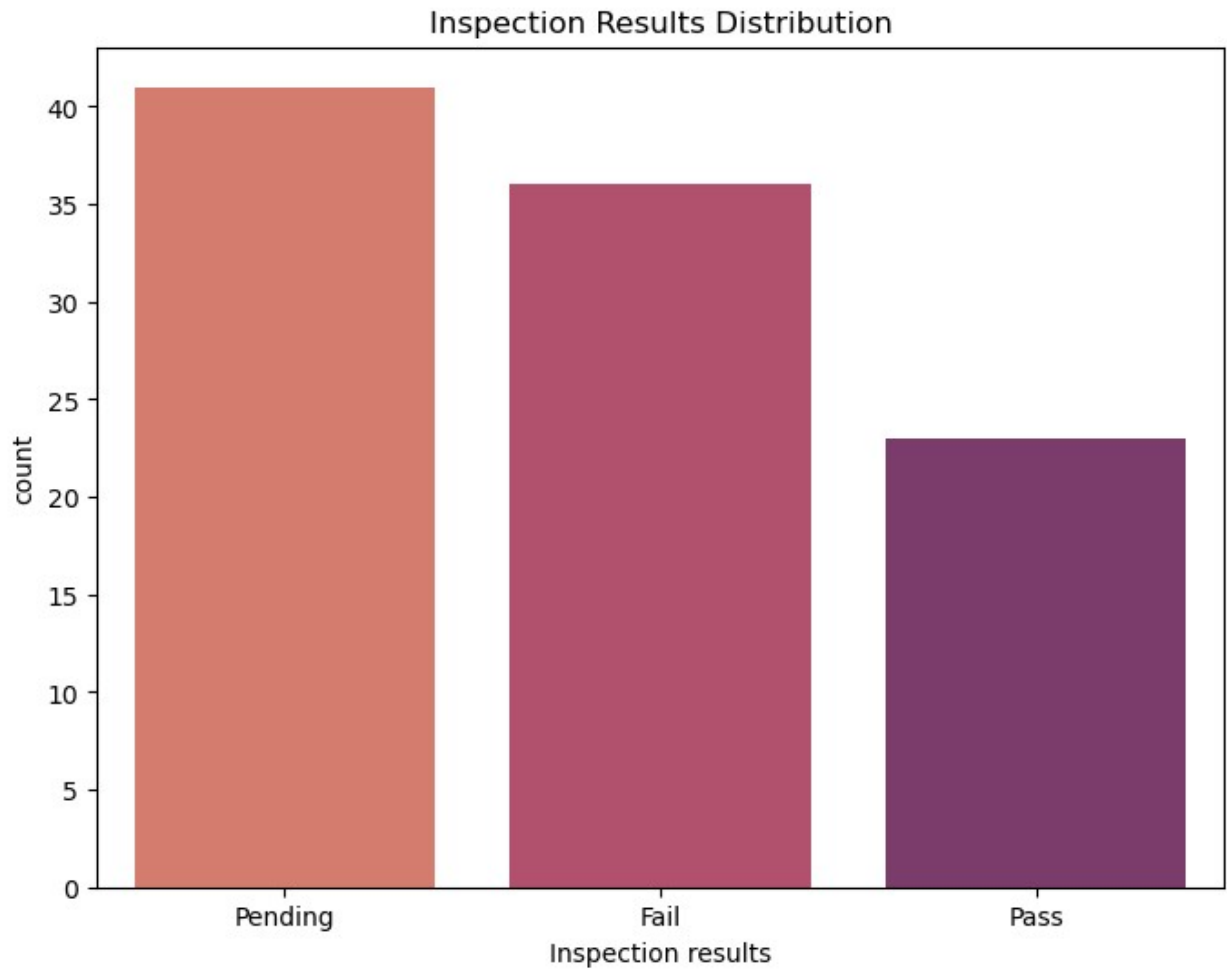


Quality Control Analysis


```
# Defect rates by product type and transportation mode
plt.figure(figsize=(14, 6))
sns.barplot(data=df, x='Product type', y='Defect
rates', hue='Transportation modes')
plt.title('Average Defect Rates by Product Type and Transportation
Mode')
plt.show()
```



```
# Inspection results
plt.figure(figsize=(8, 6))
sns.countplot(x='Inspection results', data=df, palette="flare")
plt.title('Inspection Results Distribution')
plt.show()
```



```
# Defect rates vs Manufacturing costs
px.scatter(df, x='Defect rates', y='Manufacturing costs',
           color='Inspection results', trendline="ols",
           title='Defect Rates vs Manufacturing Costs')
```



Data Preprocessing for Predictive Modeling

```
# Create a target variable - High Defect Rate (1 if defect rate >
median, else 0)
median_defect = df['Defect rates'].median()
df['High_Defect'] = (df['Defect rates'] > median_defect).astype(int)

# Select features and target
features = ['Price', 'Availability', 'Number of products sold',
'Revenue generated',
'Stock levels', 'Lead times', 'Order quantities', 'Shipping
times',
'Shipping costs', 'Lead time', 'Production volumes',
'Manufacturing lead time', 'Manufacturing costs']
target = 'High_Defect'

X = df[features]
y = df[target]

# Split data into train and test sets
X_train, X_test, y_train, y_test =
train_test_split(X,y,test_size=0.3,random_state=42,stratify=y)

# Scale numerical features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.fit_transform(X_test)
```

Predictive Modeling

```
# Initialize and train logistic regression
lr = LogisticRegression(max_iter=1000, random_state=42)
lr.fit(X_train_scaled, y_train)

# Make predictions
y_pred_lr = lr.predict(X_test_scaled)
y_prob_lr = lr.predict_proba(X_test_scaled)[: , 1]

# Evaluate model
print("Logistic Regression Performance:")
print(f"Accuracy: {accuracy_score(y_test, y_pred_lr):.4f}")
print(f"ROC AUC: {roc_auc_score(y_test, y_prob_lr):.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred_lr))

Logistic Regression Performance:
Accuracy: 0.6333
ROC AUC: 0.6622
```

| Classification Report: | | | | |
|------------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.62 | 0.67 | 0.65 | 15 |
| 1 | 0.64 | 0.60 | 0.62 | 15 |
| accuracy | | | 0.63 | 30 |
| macro avg | 0.63 | 0.63 | 0.63 | 30 |
| weighted avg | 0.63 | 0.63 | 0.63 | 30 |

2. Random Forest Classifier

```
# Initialize and train random forest
rf = RandomForestClassifier(n_estimators=200, max_depth=10,
                           random_state=42, n_jobs=-1)
rf.fit(X_train_scaled, y_train)
```

```
# Make predictions
y_pred_rf = rf.predict(X_test_scaled)
y_prob_rf = rf.predict_proba(X_test_scaled)[:, 1]
```

```
# Evaluate model
print("Random Forest Performance:")
print(f"Accuracy: {accuracy_score(y_test, y_pred_rf):.4f}")
print(f"ROC AUC: {roc_auc_score(y_test, y_prob_rf):.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred_rf))
```

Random Forest Performance:
Accuracy: 0.4667
ROC AUC: 0.5044

| Classification Report: | | | | |
|------------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.46 | 0.40 | 0.43 | 15 |
| 1 | 0.47 | 0.53 | 0.50 | 15 |
| accuracy | | | 0.47 | 30 |
| macro avg | 0.47 | 0.47 | 0.46 | 30 |
| weighted avg | 0.47 | 0.47 | 0.46 | 30 |